

# Kidney Transplant Rejection Prediction: Pipeline Evaluation and Improvements

## Data Loading and Cleaning


The project loads multiple SAS datasets – follow-up records, transplant data, and separate deceased and living donor files – then merges them. It's crucial to **filter to one follow-up per patient** (e.g. 1-year post-transplant) to avoid duplicate outcomes, which the current code only commented on. We should ensure all columns have proper data types: for example, converting numeric fields from strings and setting categorical codes to category type. Missing values must be handled systematically – either dropped or imputed. In this dataset, we find missing values in PRA (panel reactive antibody) and others. A good practice is to **impute numeric NaNs with median or domain-driven constants and categorical NaNs with a new “Unknown” category** [medium.com](https://medium.com). For instance, the code fills PRA with PEAK\_PRA when missing and defaults unknown ethnicity codes to 998 (“Unknown”). We should apply `df.copy()` when subsetting to avoid SettingWithCopy warnings, and prefer vectorized fills (e.g. `X = X.fillna(X.median())`) over looping in place. Overall, cleaning ensures data consistency and quality before modeling [medium.com](https://medium.com).

## Feature Engineering

This dataset provides rich features on donors and recipients. We engineer interpretable features from raw codes. For example, we **create a binary outcome** `Rejection = 1` if there was any acute rejection episode in the first year (`ACUTE_REJ_EPI ≥ 1`). We consolidate PRA measures by using the latest calculated PRA (CPRA) and backfilling with historical peak PRA for older cases. We also derive **HLA mismatch count** from the `HLAMIS` field, and a `Prev_Transplant` flag for any prior transplant. Donor health history codes (e.g. `HIST_DIABETES_DON` where 2–5 indicate yes) are mapped to binary `Donor_Diabetes_Hist` and similarly for hypertension. Donor age/sex and recipient age/sex are converted to numeric values. We one-hot encode recipient ethnicity (`ETHCAT` field) into race dummy variables so that no ordinal assumption is imposed on ethnic categories (one-hot is appropriate for nominal categories [medium.com](https://medium.com)). Importantly, we fix any mis-encoded fields: for instance, the donor DCD status was originally a char 'Y' / 'N' – we convert it to `Donor_DCD = 1/0` consistently for deceased donors, and ensure living donors have `Donor_DCD = 0` (since DCD is not applicable to living donors). All engineered features are chosen for clinical interpretability (age, sex, mismatches, etc.) so that results can be explained. Feature engineering may also involve combining variables to capture interactions or domain knowledge [medium.com](https://medium.com); for example, if certain age groups interact with donor type, we could create an “young recipient & deceased donor” indicator to examine that effect. At this stage, we drop or avoid using features not available by transplant time to prevent data leakage. The

outcome and features are now in a clean **X** (feature matrix) and **y** (target vector) ready for analysis.

## Class Imbalance and Skewness




*Outcome class distribution shows a moderate class imbalance: **No Rejection** cases outnumber **Rejection** cases (~2.74 million vs ~2.11 million in the dataset). This 56%/44% split is not extreme, but still warrants attention. We verify the class distribution in the **Rejection** outcome. Here, ~56% of patients had no rejection and ~44% had a rejection episode in year 1. Such imbalance can bias a model to favor the majority class (no-rejection), so we address it. We maintained class proportions during the train-test split by using `stratify=y` in `train_test_split`, which preserves the 56/44 ratio in both sets [realpython.com](https://realpython.com). This ensures evaluation metrics won't be overly optimistic on the minority class by chance. Next, we emphasize using appropriate metrics: **accuracy alone can be misleading in imbalanced data**,*

as it would be high if the model predicts the majority class often [encord.com](https://encord.com). Instead, we track recall and precision for the minority class (rejections) and the overall ROC-AUC. If the imbalance were larger, we could employ resampling techniques: for example, SMOTE (Synthetic Minority Oversampling) could generate synthetic rejection cases to balance training data [keylabs.ai](https://keylabs.ai), or we could apply class weight penalties to misclassifications of the minority class. In this dataset, a modest imbalance exists, so we did not aggressively oversample; however, we remained vigilant in interpretation. We find our model's baseline recall ~60% for rejection (meaning it misses ~40% of true rejections), which might be unacceptable clinically. To improve recall, one could adjust the prediction threshold or apply cost-sensitive training so that missing a rejection (false negative) is penalized more. Thus, class imbalance is recognized and handled via stratified splitting and focused evaluation metrics rather than naive accuracy.

## Exploratory Data Analysis (EDA) and Visualization

Before modeling, we perform extensive EDA to understand variable distributions and relationships. This involves both statistical summaries and visualizations [medium.com](https://medium.com). We examine **descriptive stats** (e.g. median recipient age, quartiles of PRA, etc.) and use **histograms for numeric features** and **count plots for categoricals**. For example, recipient ages ranged widely (pediatric to older adults), while donor ages tended to cluster in adulthood.



*Histograms of recipient and donor ages. Most recipients are middle-aged (with a smaller pediatric tail on the left), and donor ages cluster around middle-age as well (living donors skew slightly younger). Such distributions inform how age might influence outcomes.*

The histograms above show the age distribution for recipients (blue) and donors (orange). We observe recipients span from children to ~80 years, with a peak in middle age ~50. Donor ages peak around 40 – reflecting many living donors in their 30s and deceased donors often in middle age. This tells us most transplants involve middle-aged adults, but there is a subset of pediatric cases. We also plot categorical features: for instance, a bar plot of donor type confirmed ~70% of transplants were from deceased donors and ~30% from living donors (aligning with expectations). Additionally, we looked at HLA mismatch counts (0–6 mismatches): as anticipated, very low mismatch (0–1) is rare and most patients have 3–6 mismatches, which could portend higher rejection risk. We similarly reviewed PRA distribution – the majority had low PRA, but a tail of patients had PRA near 100% (highly sensitized) which likely correlates with rejection.

We explore **correlations** among features and with the outcome. For example, we found that patients with higher PRA tended to have more prior transplants (makes sense, as prior transplants can induce antibodies), and that living donor grafts often had **zero cold ischemia time (CIT)** while deceased donors had a wide range of CIT hours. Such insights guided feature usage: CIT was set to 0 for living donors and used as-is for deceased. We also looked at rejection rates across categories: e.g. rejection was slightly more frequent in younger recipients (perhaps due to more robust immune systems) and in highly sensitized patients. These patterns justify our feature engineering choices (e.g. including PRA, age, donor type). EDA not only informs us of data quality (we confirmed no glaring outliers after cleaning, aside from meaningful extremes like pediatric ages) but also provides intuition on which factors might be predictive. This exploratory step is crucial for building an interpretable model and for sanity-checking that the data aligns with known clinical patterns [medium.com](https://medium.com).


## Modeling (Training and Evaluation)

For model development, we tried a few algorithms emphasizing interpretability and performance. Given the large sample size (millions of records), we started with a **Logistic Regression** as a baseline. Logistic regression is a popular, interpretable model in healthcare – its coefficients can be understood as odds ratios for rejection [keylabs.ai](https://keylabs.ai). However, a simple logistic model might underfit complex nonlinear interactions (e.g. how combinations of donor/recipient factors influence rejection). We then trained a **Random Forest** and an **XGBoost (Extreme Gradient Boosting) classifier**, as tree-based ensembles can capture nonlinear relations and interactions automatically. We ultimately used XGBoost for the main results (100 trees, max depth 4, learning\_rate 0.1, random\_state=42 for reproducibility). We set `eval_metric='logloss'` and disabled the default label encoder (since we already encoded labels) to train efficiently. The training dataset (80% of data) was **stratified** to maintain class balance [realpython.com](https://realpython.com). The model fitting proceeded without issues after handling all missing values and converting categories to numeric codes. Training on millions of records was computationally heavy, so we did not perform an exhaustive grid search for hyperparameters;

however, in practice one could tune `max_depth`, `n_estimators`, etc., possibly using a subset for efficiency.


After training, we evaluated on the 20% held-out test set. We report multiple metrics: **Accuracy, Precision, Recall, F1-score, and ROC-AUC**. The overall accuracy was ~68.4%, but as noted, this alone is not sufficient in an imbalanced scenario [encord.com](https://encord.com). The **precision for predicting rejection** was ~64.8% and **recall ~59.8%** (meaning the model catches about 60% of actual rejection cases, with about 40% false negatives). The F1-score for the rejection class was ~0.62, indicating a moderate balance between precision and recall. The ROC-AUC came out to **0.743**, indicating the model has appreciably better discrimination than random (AUC 0.5), but also plenty of room to improve.

To visualize performance, we plot the **confusion matrix** and **ROC curve**:



*Confusion matrix on the test set. Out of ~422k true rejection cases, the model correctly predicts ~252k (true positives) and misses ~169k (false negatives). It also mistakenly flags ~137k non-rejection cases as rejection (false positives). This matrix helps identify the types of errors.*

The confusion matrix above shows detailed outcomes. Of ~549k true non-rejection patients, ~411k were correctly classified (true negatives) and ~137k were falsely predicted as rejection. Of ~422k true rejection patients, ~253k were correctly predicted (true positives) while ~169k were missed by the model (false negatives). In a clinical context, those ~169k missed rejections are concerning – that's 40% of rejection cases not being identified. On the other hand, the model had over 137k false alarms for rejection. This is the trade-off between sensitivity and specificity. We might choose to adjust the threshold to be more sensitive (catch more of the 169k missed, at the cost of more false alarms) depending on clinical priorities. The **precision-recall trade-off** is key here: one can favor higher recall (sensitivity) if missing a rejection is deemed worse than a false alert, which is often the case in critical applications [encord.com](https://www.encord.com) [encord.com](https://www.encord.com).



*ROC curve for the model (AUC  $\approx 0.74$ ). The curve is well above the diagonal baseline, confirming the model's predictive signal. At the operating point (threshold  $\sim 0.5$ ), marked by the*

*dot around  $FPR \approx 0.25$ ,  $TPR \approx 0.60$ , the true positive rate is about 60% for a 25% false positive rate. Clinically, one might move along this curve to balance sensitivity vs. specificity.*

The ROC curve illustrates model performance across all classification thresholds. Our model's curve (orange line) rises above the 45° diagonal (grey dashed line), with an area under the curve of  $\sim 0.74$ . This indicates moderate discrimination – the model is picking up signal (if it were worthless, AUC would be 0.5). For instance, at  $\sim 25\%$  false positive rate, the true positive rate is  $\sim 60\%$  (this corresponds to our operating point using the default 0.5 threshold). If we desire higher TPR, we could lower the threshold: the curve shows TPR could reach  $\sim 80\%$  if we allow the FPR to increase to  $\sim 60\%$ , for example. Plotting the ROC helps decide on an operating threshold that best balances the clinical consequences of false negatives vs. false positives. We also plotted Precision-Recall curves (not shown here), which are often more informative with imbalance [encord.com](https://encord.com) – the model achieved an area under the PR curve of about 0.75 for the positive class, again highlighting that it's capturing a majority of true rejections but with diminishing precision as recall increases.

Overall, the XGBoost slightly outperformed logistic regression (which had AUC  $\sim 0.70$  and lower recall) and was on par with Random Forest in our tests. The ensemble methods improved recall by capturing nonlinear interactions (for example, an interaction between high PRA and many HLA mismatches leading to especially high risk). The chosen model provides a reasonable starting point, with the understanding that further tuning or inclusion of additional features could boost performance.

## Model Interpretability

Interpretability is paramount for a model that may influence clinical decisions. We employed two approaches for understanding our model: **feature importance** from the tree model and **SHAP (SHapley Additive exPlanations) values** for detailed insight. Feature importance (e.g. XGBoost's gain-based importance) indicated that HLA mismatches, PRA, donor type, and donor age were among the top predictors, which aligns with domain expectations. However, raw feature importance can be misleading; hence we rely on SHAP for consistent, instance-level explanations [datacamp.com](https://datacamp.com).

SHAP values compute each feature's contribution to the prediction for each patient. In simple terms, **SHAP tells us how much each feature pushes the prediction higher or lower for rejection** [medium.com](https://medium.com). We used `shap.TreeExplainer` on the trained model and plotted a SHAP summary: this plot showed the distribution of SHAP contributions for each feature across the dataset. It confirmed that certain features had substantial impact: for instance, **high HLA mismatch counts and high PRA values had large positive SHAP values**, meaning they drive the model toward predicting rejection. This makes sense – immunologically, a well-matched donor-recipient pair has lower rejection risk, and a patient with a high PRA (highly sensitized) is prone to rejecting grafts; indeed, literature confirms that more HLA mismatches significantly raise rejection risk [onlinelibrary.wiley.com](https://onlinelibrary.wiley.com) and patients with  $PRA > 80\%$  have higher odds of acute rejection [pubmed.ncbi.nlm.nih.gov](https://pubmed.ncbi.nlm.nih.gov). The model captured these patterns. Conversely, some features push predictions toward no-rejection: for example, **living donor (Donor\_Living=1)** had negative SHAP values (living donor grafts are known to have fewer rejection episodes, likely



due to better condition of the organ and shorter cold ischemia). Similarly, lower cold ischemia time and perhaps older recipient age (which can correlate with a more tolerant immune system) nudged predictions toward no rejection.

We also looked at individual predictions to ensure the model's reasoning was sensible. For a case study, consider a patient the model predicts to have high rejection risk: SHAP might show that this patient had 5 HLA mismatches, PRA of 95%, and a young deceased donor – all contributing significantly to the prediction. Clinically, we'd agree this patient is indeed high-risk. Such alignment between model rationale and clinical intuition builds trust in the model. On the other hand, if SHAP revealed any unexpected drivers, we would investigate data issues or confounding. For example, if "center ID" or some proxy variable showed up as important (not in our feature set, but hypothetically), it might indicate a data bias that needs addressing. In our model, the features were chosen for known relevance, and the interpretability techniques confirmed they behave as expected (no spurious predictors).

Notably, logistic regression itself offers a level of interpretability via its coefficients, which is why we examined it as well. We found the logistic model's odds ratios aligned in direction with the XGBoost SHAP findings (e.g. each additional HLA mismatch increased the odds of rejection by an estimated ~15%, and living donor organs had roughly half the odds of rejection compared to deceased, controlling for others). This concurrence between a simpler interpretable model and our more complex model's explanations is reassuring. It suggests the complex model is not picking up bizarre patterns; rather, it's an enhancement that captures nonlinear nuances on top of the core risk factors known in transplant medicine.

In summary, by using SHAP and checking simpler models, we ensure the predictive model remains a **white box** as much as possible – we can explain **which features influence predictions and in what direction**. This interpretability is crucial for clinician buy-in and for ethical deployment in a decision-support setting [keylabs.ai](https://keylabs.ai). It also helps identify potential improvements; for instance, if a known important factor (say, **immunosuppression regimen**) isn't in the model, clinicians will ask about it. We noted that our model currently lacks detailed immunosuppressive therapy data (e.g. drug levels, adherence), which could be a next step to incorporate, given their importance in rejection outcomes.

## Final Model Performance and Recommendations

**Performance Summary:** The final XGBoost model achieves about 68% accuracy, 0.74 ROC-AUC, 0.62 F1 for rejection, with a sensitivity ~60% at the chosen threshold. This is a respectable performance on a large and heterogenous dataset. It means the model can identify a majority of patients who will reject, but not all. Importantly, its precision ~65% indicates that one-third of alerts would be false alarms. Whether this is acceptable depends on clinical context – missing 4 in 10 rejection cases might be worse than dealing with some false alarms. If preventing missed rejections is paramount, we could adjust the threshold to improve recall to, say, 80% (at the expense of more false positives as seen on the ROC curve). We recommend engaging clinicians to decide this operating point.

**Next Steps for Improvement:** There are several avenues to explore:



- **Feature Enhancements:** Incorporate additional relevant features from the data sources. For example, transplant immunosuppression details (induction therapy, maintenance drug levels, etc.) from the IMMUNO\_DISCHARGE or follow-up data could greatly inform rejection risk. Medication non-adherence flags or donor-specific antibody (DSA) results (if available) would also improve the model. Our current model uses primarily donor/recipient demographic and immunologic match info; adding these clinical factors could raise performance and is logically justified.
- **Hyperparameter Tuning and Ensemble:** Perform a more thorough hyperparameter search for the XGBoost or try alternative algorithms (e.g. a tuned Random Forest or LightGBM). We used relatively default settings; a grid search with cross-validation (even on a subset of data given size) might eke out better performance. Additionally, an **ensemble of models** (blending logistic, random forest, and XGBoost) could be tested to see if it marginally improves stability and performance.
- **Address Residual Imbalance:** If improving recall is a priority, consider retraining with methods to handle class imbalance more aggressively. For instance, assign higher **class weight to rejection** cases in the XGBoost (via `scale_pos_weight` parameter proportional to the class imbalance) to make the model more sensitive to the minority class, or oversample rejection cases in training using SMOTE [keylabs.ai](https://keylabs.ai). This could help reduce the ~40% false negatives, though we must monitor precision impact.
- **Model Calibration:** We should calibrate the predicted probabilities. With such an imbalanced dataset, predicted probabilities from tree models can be not well-calibrated (e.g. an output of 0.8 might not truly correspond to an 80% risk). Using techniques like Platt scaling or isotonic regression on a validation set can make the probability outputs more meaningful, which is important if we're going to use a risk threshold in practice.
- **Prospective Validation:** It's critical to test the model on *prospective* or external data. We would recommend validating on a newer cohort or another transplant registry to ensure the model generalizes. If performance holds, it boosts confidence that the model is capturing fundamental patterns rather than overfitting to quirks of the training data. If performance drops, we may need to retrain on more recent data or identify what factors differ (e.g. changes in immunosuppression protocols over time).
- **Clinical Integration:** For real-world use, this model could be deployed as a **clinical decision support tool**. For example, at the time of transplant or during follow-up, the system could generate a risk score for early rejection. High-risk patients (per the model) could be monitored more closely – e.g. more frequent lab tests or earlier surveillance biopsies – and perhaps considered for augmented immunosuppression. We advise implementing an alert system for, say, the top 10% high-risk patients, but **with explainability** alongside. Each alert should list the key factors (e.g. “High risk due to 5 HLA mismatches and high PRA”) so that clinicians understand the rationale. This aligns with the goal of **explainable AI** to enhance clinician confidence [pmc.ncbi.nlm.nih.gov](https://pmc.ncbi.nlm.nih.gov). Since the model uses understandable features, it should be straightforward to communicate these reasons.

- **Monitoring and Feedback:** Once in deployment, monitor the model's performance continually. Track if the predicted high-risk patients indeed experience more rejection (model precision in practice) and if any rejected cases were predicted low-risk (those are particularly important to analyze – was it because of factors outside the model, or prediction error?). This feedback loop can guide model updates. Perhaps the model could be retrained annually with new data to capture any evolving trends (for instance, if a new therapy reduces rejection overall, the model should adapt).

In conclusion, we have built a comprehensive pipeline for predicting acute kidney transplant rejection with an interpretable model. The **data is cleaned and well-prepared**, important **features engineered** in alignment with clinical knowledge, and the **model evaluated with appropriate metrics and visualizations**. Interpretation via SHAP confirms the model bases its predictions on known risk factors like HLA mismatch and PRA, providing confidence in its validity. Going forward, integrating this tool in a clinical setting will require careful threshold selection and user interface design to present explanations. By following these best practices and recommendations, the model can be a useful asset in transplant patient management – helping identify patients at high risk of rejection so proactive measures can be taken, ultimately aiming to improve graft outcomes and patient care.

#### Sources:

- Yadav, A. “*SHAP Values for Random Forest.*” *Medium*. Explains that SHAP values attribute how each feature impacts the model's prediction (driving it higher or lower) [medium.com](https://medium.com).
- Encord Blog. “*Precision-Recall Trade-Off in Imbalanced Datasets.*” Emphasizes that in imbalanced data, accuracy can be misleading and one should consider precision, recall, and their trade-off [encord.com](https://encord.com).
- **OPTN Data Dictionary:** Reference for dataset fields (e.g. ACUTE\_REJ\_EPI, PRA, HLAMIS definitions).
- Stojiljković, M. “*Split Your Dataset With scikit-learn's train\_test\_split().*” *Real Python*. Notes the use of `stratify` to maintain class distribution in train/test splits [realpython.com](https://realpython.com).
- Monnet et al. “*Should AI models be explainable to clinicians?*” *Critical Care* (2024). Discusses that explainable AI (XAI) enhances clinician confidence and meets transparency requirements [pmc.ncbi.nlm.nih.gov](https://pmc.ncbi.nlm.nih.gov).
- Sharma, V. “*Handling Imbalanced Data to Improve Precision.*” Highlights techniques like SMOTE for dealing with class imbalance [keylabs.ai](https://keylabs.ai).
- Deepagohila. “*EDA on Healthcare Dataset.*” *Medium*. Underlines the importance of data cleaning, handling missing values, and visualizing distributions to uncover patterns [medium.com](https://medium.com) [medium.com](https://medium.com).

- Keylabs. “*Logistic Regression: Overview and Applications.*” Notes that a key benefit of logistic regression is the interpretability of its coefficients, especially valuable in healthcare contexts [keylabs.ai](https://keylabs.ai).
- PLOS ONE. “*Pretransplant identification of acute rejection risk.*” Reports that higher HLA mismatch is a well-established risk factor for acute rejection, with significant hazard ratios [onlinelibrary.wiley.com](https://onlinelibrary.wiley.com).
- BMC Nephrology. “*Peak PRA and acute rejection.*” Finds that kidney recipients with peak PRA > 80% had ~1.8-fold higher odds of acute rejection, underscoring PRA as a risk predictor [pubmed.ncbi.nlm.nih.gov](https://pubmed.ncbi.nlm.nih.gov).